

## GPU

```
!nvidia-smi
```

```
Sun Nov 30 02:04:26 2025
```

NVIDIA-SMI 550.54.15			Driver Version: 550.54.15		CUDA Version: 12.4		
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr.	ECC
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.	MIG M.
0	Tesla T4	Off	00000000:00:04.0	Off	0		
N/A	49C	P0	27W / 70W	376MiB / 15360MiB	0%	Default	N/A

Processes:				
GPU	GI	CI	PID	Type
ID	ID			Process name

## 1. Library Imports

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
import seaborn as sns
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor
import joblib
from sklearn.inspection import PartialDependenceDisplay
from tensorflow import keras
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score
```

```
import numpy as np
from tensorflow.keras.callbacks import EarlyStopping
import tensorflow as tf
```

## 2. Data Loading

```
df = pd.read_csv('/content/Superstore.csv')
```

## 3. Initial Data Overview

```
df.head()
```

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Country	City	...	Postal Code	Region	Product ID	Category	Sub-Category
0	1	CA-2016-152156	11/8/2016	11/11/2016	Second Class	CG-12520	Claire Gute	Consumer	United States	Henderson	...	42420	South	FUR-BO-10001798	Furniture	Bookcases
1	2	CA-2016-152156	11/8/2016	11/11/2016	Second Class	CG-12520	Claire Gute	Consumer	United States	Henderson	...	42420	South	FUR-CH-10000454	Furniture	Chairs
2	3	CA-2016-138688	6/12/2016	6/16/2016	Second Class	DV-13045	Darrin Van Huff	Corporate	United States	Los Angeles	...	90036	West	OFF-LA-10000240	Office Supplies	Labels
3	4	US-2015-108966	10/11/2015	10/18/2015	Standard Class	SO-20335	Sean O'Donnell	Consumer	United States	Fort Lauderdale	...	33311	South	FUR-TA-10000577	Furniture	Tables
4	5	US-2015-108966	10/11/2015	10/18/2015	Standard Class	SO-20335	Sean O'Donnell	Consumer	United States	Fort Lauderdale	...	33311	South	OFF-ST-10000760	Office Supplies	Storage

5 rows × 21 columns

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9994 entries, 0 to 9993
Data columns (total 21 columns):
 #   Column      Non-Null Count  Dtype  
 --- 
 0   Row ID      9994 non-null   int64  
 1   Order ID    9994 non-null   object  
 2   Order Date  9994 non-null   object  
 3   Ship Date   9994 non-null   object  
 4   Ship Mode   9994 non-null   object  
 5   Customer ID 9994 non-null   object  
 6   Customer Name 9994 non-null   object  
 7   Segment     9994 non-null   object  
 8   Country     9994 non-null   object  
 9   City        9994 non-null   object  
 10  State       9994 non-null   object  
 11  Postal Code 9994 non-null   int64  
 12  Region      9994 non-null   object  
 13  Product ID  9994 non-null   object  
 14  Category    9994 non-null   object  
 15  Sub-Category 9994 non-null   object  
 16  Product Name 9994 non-null   object  
 17  Sales        9994 non-null   float64 
 18  Quantity     9994 non-null   int64  
 19  Discount     9994 non-null   float64 
 20  Profit       9994 non-null   float64 
dtypes: float64(3), int64(3), object(15)
memory usage: 1.6+ MB
```

```
df.describe()
```

	Row ID	Postal Code	Sales	Quantity	Discount	Profit	grid icon
count	9994.000000	9994.000000	9994.000000	9994.000000	9994.000000	9994.000000	grid icon
mean	4997.500000	55190.379428	229.858001	3.789574	0.156203	28.656896	bar chart icon
std	2885.163629	32063.693350	623.245101	2.225110	0.206452	234.260108	
min	1.000000	1040.000000	0.444000	1.000000	0.000000	-6599.978000	
25%	2499.250000	23223.000000	17.280000	2.000000	0.000000	1.728750	
50%	4997.500000	56430.500000	54.490000	3.000000	0.200000	8.666500	
75%	7495.750000	90008.000000	209.940000	5.000000	0.200000	29.364000	
max	9994.000000	99301.000000	22638.480000	14.000000	0.800000	8399.976000	

```
df.isnull().sum()
```

	0
<b>Row ID</b>	0
<b>Order ID</b>	0
<b>Order Date</b>	0
<b>Ship Date</b>	0
<b>Ship Mode</b>	0
<b>Customer ID</b>	0
<b>Customer Name</b>	0
<b>Segment</b>	0
<b>Country</b>	0
<b>City</b>	0
<b>State</b>	0
<b>Postal Code</b>	0
<b>Region</b>	0
<b>Product ID</b>	0
<b>Category</b>	0
<b>Sub-Category</b>	0
<b>Product Name</b>	0
<b>Sales</b>	0
<b>Quantity</b>	0
<b>Discount</b>	0
<b>Profit</b>	0


**dtype:** int64

#### 4. Data Visualization and Exploration


```
# Sales, Profit, Quantity distribution
plt.figure(figsize=(16,4))
plt.subplot(1,3,1)
sns.histplot(df['Sales'], bins=50, kde=True)
plt.title("Sales Distribution")
```

```
plt.subplot(1,3,2)
sns.histplot(df['Profit'], bins=50, kde=True, color='green')
plt.title("Profit Distribution")


plt.subplot(1,3,3)
sns.histplot(df['Quantity'], bins=50, kde=True, color='orange')
plt.title("Quantity Distribution")
plt.show()
```



```
# Discount distribution
plt.figure(figsize=(6,4))
sns.countplot(x='Discount', data=df)
plt.title("Discount Counts")
plt.show()
```



```
# Sales vs Profit scatter plot
plt.figure(figsize=(6,4))
sns.scatterplot(x='Sales', y='Profit', hue='Category', data=df, alpha=0.6)
plt.title("Sales vs Profit by Category")
plt.show()
```





```
# Profit by Category and Sub-Category
plt.figure(figsize=(12,5))
sns.boxplot(x='Category', y='Profit', data=df)
plt.title("Profit Distribution by Category")
plt.show()

plt.figure(figsize=(15,5))
sns.boxplot(x='Sub-Category', y='Profit', data=df)
plt.xticks(rotation=45)
plt.title("Profit Distribution by Sub-Category")
plt.show()
```



## Profit Distribution by Category

```
# Ship Mode counts
plt.figure(figsize=(6,4))
sns.countplot(x='Ship Mode', data=df)
plt.title("Ship Mode Counts")
plt.show()
```



## 5. Feature Engineering and Preprocessing


```
df['Order Date'] = pd.to_datetime(df['Order Date'])
df['Ship Date'] = pd.to_datetime(df['Ship Date'])

df['OrderYear'] = df['Order Date'].dt.year.astype(int)
df['OrderMonth'] = df['Order Date'].dt.month.astype(int)
```


```
# Orders over Time
df['OrderYearMonth'] = df['Order Date'].dt.to_period('M')
orders_per_month = df.groupby('OrderYearMonth').size()
```

```
plt.figure(figsize=(12,4))
orders_per_month.plot()
plt.title("Number of Orders Over Time")
plt.xlabel("Year-Month")
```

```
plt.ylabel("Number of Orders")
plt.show()
```



```
# Correlation heatmap for numeric features
plt.figure(figsize=(8,6))
sns.heatmap(df[['Sales','Quantity','Discount','Profit']].corr(), annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()
```



```
df.columns
```

```
Index(['Row ID', 'Order ID', 'Order Date', 'Ship Date', 'Ship Mode',
       'Customer ID', 'Customer Name', 'Segment', 'Country', 'City', 'State',
       'Postal Code', 'Region', 'Product ID', 'Category', 'Sub-Category',
       'Product Name', 'Sales', 'Quantity', 'Discount', 'Profit', 'OrderYear',
       'OrderMonth', 'OrderYearMonth'],
      dtype='object')
```

```
df.shape
```

```
(9994, 24)
```

```
print(df.City.unique())
df.City.value_counts()
```

531

count

City

New York City	915
Los Angeles	747
Philadelphia	537
San Francisco	510
Seattle	428
...	...
Abilene	1
Montebello	1
Kissimmee	1
Danbury	1
Springdale	1

531 rows × 1 columns

dtype: int64

```
print(df.Country.nunique())
df.Country.value_counts()
```

1

count

Country

United States	9994
---------------	------

dtype: int64

```
print(df.State.nunique())
df.State.value_counts()
```



49

State	count
California	2001
New York	1128
Texas	985
Pennsylvania	587
Washington	506
Illinois	492
Ohio	469
Florida	383
Michigan	255
North Carolina	249
Arizona	224
Virginia	224
Georgia	184
Tennessee	183
Colorado	182
Indiana	149
Kentucky	139
Massachusetts	135
New Jersey	130
Oregon	124

## 5.1. Dropping Unnecessary Columns

Wisconsin 110  
 Maryland 105

```
# Drop UnNeeded Columns
cols_to_drop = [
    'Row ID', 'Order ID', 'Customer ID', 'Customer Name', 'Product ID', 'Country'
]
df = df.drop(columns=cols_to_drop)
```

df.OrderYearMonth[0]

Alabama 61

```
Period('2016-11', 'M') 60
Arkansas
```

✓ 5.2. Date Feature Extraction and Dropping Date Columns

Rhode Island	56
Utah	53

```
# Extract New Features From Dates
df['OrderYear'] = df['Order Date'].dt.year
df['OrderMonth'] = df['Order Date'].dt.month
df['ShippingDelay'] = (df['Ship Date'] - df['Order Date']).dt.days

df = df.drop(columns=['Order Date', 'Ship Date', 'OrderYearMonth'])
```

NEBRASKA 50

✓ 5.3. Categorical Feature Encoding

New Mexico	37
Iowa	30

```
# Label Encoding
cat_cols = ['Ship Mode', 'Segment', 'Region', 'Category', 'Sub-Category', 'Product Name', "City", "State"]
encoders = {}
for col in cat_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    encoders[col] = le
```

✓ 5.4. Outlier Handling (Clipping)

Vermont	11
District of Columbia	10

```
# Handling Outliers
lower = df['Profit'].quantile(0.01)
upper = df['Profit'].quantile(0.99)
df['Profit_clipped'] = df['Profit'].clip(lower, upper)

# Sales
lower_sales = df['Sales'].quantile(0.01)
upper_sales = df['Sales'].quantile(0.99)
df['Sales_clipped'] = df['Sales'].clip(lower_sales, upper_sales)

# Quantity
lower_qty = df['Quantity'].quantile(0.01)
upper_qty = df['Quantity'].quantile(0.99)
df['Quantity_clipped'] = df['Quantity'].clip(lower_qty, upper_qty)
```

✓ 6. Data Splitting and Scaling

```
# Split Data
v = df['Sales_clipped']
```

```
y = df['Sales_Clipped']  
X = df.drop(columns=['Sales', 'Sales_Clipped', 'Profit'])  
  
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42  
)
```

```
# Save the Feature Order  
joblib.dump(X.columns.tolist(), "feature_order.joblib")  
  
['feature_order.joblib']
```

```
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)
```

## ▼ 7. ML Model Training and Evaluation

### ▼ 7.1. Support Vector Regressor (SVR)

```
svr = SVR(kernel='rbf', C=100, epsilon=0.1)  
svr.fit(X_train_scaled, y_train)  
y_pred_svr = svr.predict(X_test_scaled)  
  
# Evaluate Decision Tree  
mae_dt = mean_absolute_error(y_test, y_pred_svr)  
mse_dt = mean_squared_error(y_test, y_pred_svr)  
rmse_dt = mse_dt ** 0.5  
r2_dt = r2_score(y_test, y_pred_svr)
```

```
print("SVR Evaluation:")  
print("-----")  
print("MAE :", mae_dt)  
print("MSE :", mse_dt)  
print("RMSE:", rmse_dt)  
print("R² : ", r2_dt)
```

```
SVR Evaluation:  
-----  
MAE : 96.98692412337483  
MSE : 62853.12599573185  
RMSE: 250.7052572159823  
R² : 0.6508089961379162
```

## 7.2. Decision Tree Regressor

```
# Train Decision Tree
dt = DecisionTreeRegressor(
    max_depth=10,
    min_samples_leaf=5,
    random_state=42
)

dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_test)
```

```
# Evaluate Decision Tree
mae_dt = mean_absolute_error(y_test, y_pred_dt)
mse_dt = mean_squared_error(y_test, y_pred_dt)
rmse_dt = mse_dt ** 0.5
r2_dt = r2_score(y_test, y_pred_dt)

print("Decision Tree Evaluation:")
print("-----")
print("MAE :", mae_dt)
print("MSE :", mse_dt)
print("RMSE:", rmse_dt)
print("R²  :", r2_dt)
```

Decision Tree Evaluation:

-----  
MAE : 72.54262425810613  
MSE : 38164.71877835814  
RMSE: 195.3579247902632  
R² : 0.7879695520118764

## 7.3. Random Forest Regressor

```
# Train Random Forest model
rf = RandomForestRegressor(n_estimators=500, max_depth=15, min_samples_leaf=5,
                           random_state=42, n_jobs=-1)

rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
```

```
# Evaluate Random Forest Model
mae_rf = mean_absolute_error(y_test, y_pred_rf)
mse_rf = mean_squared_error(y_test, y_pred_rf)
rmse_rf = mse_rf ** 0.5
r2_rf = r2_score(y_test, y_pred_rf)
```

```
print("Random Forest Evaluation:")
print("-----")
print("MAE :", mae_rf)
print("MSE :", mse_rf)
print("RMSE:", rmse_rf)
print("R²  :", r2_rf)

Random Forest Evaluation:
-----
MAE : 64.06244436073729
MSE : 31078.00637819716
RMSE: 176.289552663217
R²  : 0.8273409623894944
```

## ▼ 8. Model Persistence (Saving Models)

```
# Save SVR model
joblib.dump(svr, 'SVR_model.joblib')
print("SVR model saved as SVR_model.joblib")

# Save Decision Tree model
joblib.dump(dt, 'DecisionTree_model.joblib')
print("Decision Tree model saved as DecisionTree_model.joblib")

# Save Random Forest model
joblib.dump(rf, 'RandomForest_model.joblib')
print("Random Forest model saved as RandomForest_model.joblib")

SVR model saved as SVR_model.joblib
Decision Tree model saved as DecisionTree_model.joblib
Random Forest model saved as RandomForest_model.joblib
```

```
# Save LabelEncoder
joblib.dump(encoders, "label_encoders.joblib")
print("LabelEncoder saved as label_encoder.joblib")

# Save the StandardScaler used for scikit-learn models
joblib.dump(scaler, 'standard_scaler_sklearn.joblib')
print("StandardScaler (for scikit-learn models) saved as standard_scaler_sklearn.joblib")


LabelEncoder saved as label_encoder.joblib
StandardScaler (for scikit-learn models) saved as standard_scaler_sklearn.joblib
```

## ▼ 9. Random Forest Model Analysis

```
# Feature Importance
```

```
importances = rf.feature_importances_
feature_names = X_train.columns
feat_imp = pd.Series(importances, index=feature_names).sort_values(ascending=False)

plt.figure(figsize=(12,6))
sns.barplot(x=feat_imp[:10], y=feat_imp.index[:10])
plt.title("Top 10 Feature Importances")
plt.xlabel("Importance")
plt.ylabel("Feature")
plt.show()
```



## 9.1. Partial Dependence Plots

```
# Top 5 Features
top_features = feat_imp.index[:5]

# Plot Partial Dependence
fig, ax = plt.subplots(figsize=(16,12))
display = PartialDependenceDisplay.from_estimator(
```


```
rf,
X_train,
features=top_features,
kind="average",
grid_resolution=50,
ax=ax
)
plt.suptitle("Partial Dependence Plots (Top 5 Features)", fontsize=16)
plt.subplots_adjust(top=0.92)
plt.show()
```



## Partial Dependence Plots (Top 5 Features)


## 9.2. Random Forest Residual Analysis

```
plt.figure(figsize=(8,6))
plt.scatter(y_test, y_pred_rf, alpha=0.5)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
plt.xlabel("Actual Sales")
plt.ylabel("Predicted Sales")
plt.title("Actual vs Predicted Sales")
plt.show()
```




```
residuals = y_test - y_pred_rf
plt.figure(figsize=(8,6))
sns.histplot(residuals, bins=50, kde=True)
plt.title("Residuals Distribution")
plt.xlabel("Residual")
```

```
plt.ylabel('Frequency')  
plt.show()
```



```
plt.figure(figsize=(8,6))  
plt.scatter(y_pred_rf, residuals, alpha=0.5)  
plt.axhline(0, color='red', linestyle='--')  
plt.xlabel("Predicted Sales")  
plt.ylabel("Residuals")  
plt.title("Residuals vs Predicted Sales")  
plt.show()
```



## ▼ 10. Neural Network Model (Keras)

```
kerasmodel = keras.models.Sequential([
    keras.layers.Input(shape=(16,)),
    keras.layers.Dense(64, activation='tanh'),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dropout(0.3),
    keras.layers.Dense(64, activation='relu'),
    keras.layers.Dropout(0.1),
    keras.layers.Dense(32, activation='relu'),
    keras.layers.Dense(1, activation='linear')
])
```

```
MyOptimizer = tf.keras.optimizers.AdamW(
    learning_rate=0.001,
    weight_decay=0.004,
    beta_1=0.9,
```

```
        beta_2=0.999,
        epsilon=1e-07,
        amsgrad=False,
        clipnorm=None,
        clipvalue=None,
        global_clipnorm=None,
        use_ema=False,
        ema_momentum=0.99,
        ema_overwrite_frequency=None,
        name="AdamW"
    )
```

```
scaler_x = StandardScaler()
scaler_y = StandardScaler()

X_train_scaled = scaler_x.fit_transform(X_train)
X_test_scaled = scaler_x.transform(X_test)

y_train_scaled = scaler_y.fit_transform(y_train.values.reshape(-1,1))
y_test_scaled = scaler_y.transform(y_test.values.reshape(-1,1))
```

```
kerasmodel.compile(
    optimizer='adam',
    loss='mse',
    metrics=['mae', 'mse']
)

early_stop = EarlyStopping(
    monitor='val_loss',
    patience=10,
    mode='min',
    restore_best_weights=True
)

history = kerasmodel.fit(
    X_train_scaled, y_train_scaled,
    validation_data=(X_test_scaled, y_test_scaled),
    epochs=100,
    batch_size=64,
    verbose=1
)
```

```
Epoch 75/100
125/125      0s 4ms/step - loss: 0.0912 - mae: 0.1645 - mse: 0.0912 - val_loss: 0.2420 - val_mae: 0.2017 - val_mse: 0.2420
Epoch 76/100
125/125      0s 4ms/step - loss: 0.0974 - mae: 0.1706 - mse: 0.0974 - val_loss: 0.2474 - val_mae: 0.2066 - val_mse: 0.2474
Epoch 77/100
125/125      0s 4ms/step - loss: 0.0986 - mae: 0.1710 - mse: 0.0986 - val_loss: 0.2307 - val_mae: 0.1932 - val_mse: 0.2307
Epoch 78/100
125/125      0s 4ms/step - loss: 0.0962 - mae: 0.1650 - mse: 0.0962 - val_loss: 0.2387 - val_mae: 0.2124 - val_mse: 0.2387
Epoch 79/100
125/125      0s 4ms/step - loss: 0.0891 - mae: 0.1691 - mse: 0.0891 - val_loss: 0.2536 - val_mae: 0.2040 - val_mse: 0.2536
Epoch 80/100
125/125      0s 4ms/step - loss: 0.0884 - mae: 0.1635 - mse: 0.0884 - val_loss: 0.2607 - val_mae: 0.2107 - val_mse: 0.2607
Epoch 81/100
125/125      0s 4ms/step - loss: 0.0936 - mae: 0.1679 - mse: 0.0936 - val_loss: 0.2364 - val_mae: 0.2035 - val_mse: 0.2364
Epoch 82/100
125/125      1s 5ms/step - loss: 0.0840 - mae: 0.1617 - mse: 0.0840 - val_loss: 0.2523 - val_mae: 0.2136 - val_mse: 0.2523
Epoch 83/100
125/125      1s 5ms/step - loss: 0.1007 - mae: 0.1740 - mse: 0.1007 - val_loss: 0.2493 - val_mae: 0.2064 - val_mse: 0.2493
Epoch 84/100
125/125      1s 4ms/step - loss: 0.0968 - mae: 0.1658 - mse: 0.0968 - val_loss: 0.2698 - val_mae: 0.2146 - val_mse: 0.2698
Epoch 85/100
125/125      0s 4ms/step - loss: 0.0833 - mae: 0.1623 - mse: 0.0833 - val_loss: 0.2629 - val_mae: 0.2060 - val_mse: 0.2629
Epoch 86/100
125/125      0s 4ms/step - loss: 0.0888 - mae: 0.1631 - mse: 0.0888 - val_loss: 0.2532 - val_mae: 0.2091 - val_mse: 0.2532
Epoch 87/100
125/125      0s 4ms/step - loss: 0.0996 - mae: 0.1690 - mse: 0.0996 - val_loss: 0.2404 - val_mae: 0.1985 - val_mse: 0.2404
Epoch 88/100
125/125      0s 4ms/step - loss: 0.0927 - mae: 0.1636 - mse: 0.0927 - val_loss: 0.2459 - val_mae: 0.2052 - val_mse: 0.2459
Epoch 89/100
125/125      1s 4ms/step - loss: 0.0862 - mae: 0.1626 - mse: 0.0862 - val_loss: 0.2412 - val_mae: 0.2016 - val_mse: 0.2412
Epoch 90/100
125/125      0s 4ms/step - loss: 0.0840 - mae: 0.1618 - mse: 0.0840 - val_loss: 0.2555 - val_mae: 0.2070 - val_mse: 0.2555
Epoch 91/100
125/125      0s 4ms/step - loss: 0.0840 - mae: 0.1576 - mse: 0.0840 - val_loss: 0.2798 - val_mae: 0.2177 - val_mse: 0.2798
Epoch 92/100
125/125      1s 4ms/step - loss: 0.0865 - mae: 0.1650 - mse: 0.0865 - val_loss: 0.2362 - val_mae: 0.1982 - val_mse: 0.2362
Epoch 93/100
125/125      0s 4ms/step - loss: 0.0819 - mae: 0.1573 - mse: 0.0819 - val_loss: 0.2379 - val_mae: 0.1963 - val_mse: 0.2379
Epoch 94/100
125/125      0s 3ms/step - loss: 0.0910 - mae: 0.1634 - mse: 0.0910 - val_loss: 0.2526 - val_mae: 0.2126 - val_mse: 0.2526
Epoch 95/100
125/125      0s 4ms/step - loss: 0.0981 - mae: 0.1664 - mse: 0.0981 - val_loss: 0.2320 - val_mae: 0.1985 - val_mse: 0.2320
Epoch 96/100
125/125      1s 4ms/step - loss: 0.0772 - mae: 0.1556 - mse: 0.0772 - val_loss: 0.2622 - val_mae: 0.2042 - val_mse: 0.2622
Epoch 97/100
125/125      0s 4ms/step - loss: 0.0805 - mae: 0.1565 - mse: 0.0805 - val_loss: 0.2346 - val_mae: 0.2015 - val_mse: 0.2346
Epoch 98/100
125/125      0s 3ms/step - loss: 0.0820 - mae: 0.1581 - mse: 0.0820 - val_loss: 0.2601 - val_mae: 0.2087 - val_mse: 0.2601
Epoch 99/100
125/125      0s 4ms/step - loss: 0.0767 - mae: 0.1556 - mse: 0.0767 - val_loss: 0.2336 - val_mae: 0.2015 - val_mse: 0.2336
Epoch 100/100
125/125      0s 3ms/step - loss: 0.0791 - mae: 0.1575 - mse: 0.0791 - val_loss: 0.2179 - val_mae: 0.1977 - val_mse: 0.2179
```

```
print(kerasmodel.summary())
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	1,088
dense_1 (Dense)	(None, 128)	8,320
dropout (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 64)	8,256
dropout_1 (Dropout)	(None, 64)	0
dense_3 (Dense)	(None, 32)	2,080
dense_4 (Dense)	(None, 1)	33

```
Total params: 59,333 (231.77 KB)
```

```
Trainable params: 19,777 (77.25 KB)
```

```
Non-trainable params: 0 (0.00 B)
```

```
Optimizer params: 39,556 (154.52 KB)
```

```
kerasmodel.save("KerasModel.keras")
```

```
# Save the StandardScaler for X (Keras model)
joblib.dump(scaler_x, 'standard_scaler_X_keras.joblib')
print("StandardScaler (for Keras X) saved as standard_scaler_X_keras.joblib")
```

```
# Save the StandardScaler for y (Keras model)
joblib.dump(scaler_y, 'standard_scaler_y_keras.joblib')
print("StandardScaler (for Keras y) saved as standard_scaler_y_keras.joblib")
```

```
StandardScaler (for Keras X) saved as standard_scaler_X_keras.joblib
StandardScaler (for Keras y) saved as standard_scaler_y_keras.joblib
```

```
y_pred_scaled = kerasmodel.predict(X_test_scaled)
y_pred = scaler_y.inverse_transform(y_pred_scaled)
```

```
63/63 ━━━━━━━━ 1s 5ms/step
```

```
NewKerasModel = keras.models.load_model('KerasModel.keras')
```

```
y_pred_scaled = NewKerasModel.predict(X_test_scaled)
```

```
63/63 ━━━━━━━━ 1s 5ms/step
```

```
print('Prediction Shape is {}'.format(y_pred_scaled.shape))
```

```
Prediction Shape is (1999, 1)
```

```
print('Prediction items are {}'.format(y_pred[:5]))
```



```
Prediction items are [[ 98.35385 ]  
[ 25.387032]  
[ 26.733635]  
[470.91507 ]  
[296.91318 ]]
```

```
ModelLoss = NewKerasModel.evaluate(X_test_scaled, y_test_scaled)  
print('Model Loss is {}'.format(ModelLoss))
```

```
63/63 ━━━━━━━━━━ 1s 9ms/step - loss: 0.2384 - mae: 0.2072 - mse: 0.2384  
Model Loss is [0.2179267257452011, 0.19773264229297638, 0.2179267257452011]
```


```
# --- Plot Training vs Validation Loss ---  
plt.figure(figsize=(10,5))  
plt.plot(history.history['loss'], label='Train Loss')  
plt.plot(history.history['val_loss'], label='Val Loss')  
plt.title('Training vs Validation Loss')  
plt.xlabel('Epochs')  
plt.ylabel('MSE Loss')  
plt.legend()  
plt.grid(True)  
plt.show()
```

```
# --- Plot Training vs Validation MAE ---  
plt.figure(figsize=(10,5))  
plt.plot(history.history['mae'], label='Train MAE')  
plt.plot(history.history['val_mae'], label='Val MAE')  
plt.title('Training vs Validation MAE')  
plt.xlabel('Epochs')  
plt.ylabel('MAE')  
plt.legend()  
plt.grid(True)  
plt.show()
```



```
y_pred = scaler_y.inverse_transform(y_pred_scaled)
y_test_original = scaler_y.inverse_transform(y_test_scaled)


plt.figure(figsize=(10,5))
plt.plot(y_test_original[:200], label='Actual')
plt.plot(y_pred[:200], label='Predicted')
plt.title("Actual vs Predicted (First 200 samples)")
plt.xlabel("Samples")
plt.ylabel("Target Value")
plt.legend()
plt.grid(True)
plt.show()
```



```
residuals = y_test_original - y_pred


plt.figure(figsize=(10,5))
plt.scatter(range(len(residuals)), residuals, alpha=0.5)
plt.axhline(0, color='black', linestyle='--')
plt.title("Residuals Plot")
plt.xlabel("Samples")
plt.ylabel("Residual (Actual - Predicted)")
```

```
plt.grid(True)  
plt.show()
```



```
plt.figure(figsize=(8,5))  
plt.hist(residuals, bins=40, alpha=0.7)  
plt.title("Residuals Distribution")  
plt.xlabel("Residual Value")  
plt.ylabel("Frequency")  
plt.grid(True)  
plt.show()
```

Residuals Distribution



```
y_pred_scaled = kerasmodel.predict(X_test_scaled)
y_pred = scaler_y.inverse_transform(y_pred_scaled)
y_test_original = scaler_y.inverse_transform(y_test_scaled)
```

```
63/63 ━━━━━━ 0s 2ms/step
```

```
mae_nn = mean_absolute_error(y_test_original, y_pred)
mse_nn = mean_squared_error(y_test_original, y_pred)
rmse_nn = mse_rf ** 0.5
r2_nn = r2_score(y_test_original, y_pred)

print("Neural Network Evaluation:")
print("-----")
print("MAE : ", mae_nn)
print("MSE : ", mse_nn)
print("RMSE: ", rmse_nn)
print("R² : ", r2_nn)
```

```
Neural Network Evaluation:
```

```
-----
MAE : 76.53025399058163
MSE : 32645.24775081094
RMSE: 176.289552663217
R² : 0.8186338920643946
```

- ❖ Hyperparameter Tuning:
- ❖ Experiment with Learning Rate and Batch Size

```
def create_and_compile_model(learning_rate):
    model = keras.models.Sequential([
        keras.layers.Input(shape=(16,)),
```

```
        keras.layers.Dense(64, activation='tanh'),
        keras.layers.Dense(128, activation='relu'),
        keras.layers.Dropout(0.3),
        keras.layers.Dense(64, activation='relu'),
        keras.layers.Dropout(0.1),
        keras.layers.Dense(32, activation='relu'),
        keras.layers.Dense(1, activation='linear')
    ))
optimizer = tf.keras.optimizers.Adam(learning_rate=learning_rate)
model.compile(
    optimizer=optimizer,
    loss='mse',
    metrics=['mae', 'mse']
)
return model

print("Defined create_and_compile_model function.")
```

Defined create\_and\_compile\_model function.

```
learning_rates = [0.01, 0.001, 0.0001]
batch_sizes = [32, 64, 128]
epochs = 50

results = []

for lr in learning_rates:
    for bs in batch_sizes:
        print(f"\nTraining with Learning Rate: {lr}, Batch Size: {bs}")
        model = create_and_compile_model(learning_rate=lr)

        history = model.fit(
            X_train_scaled, y_train_scaled,
            validation_data=(X_test_scaled, y_test_scaled),
            epochs=epochs,
            batch_size=bs,
            verbose=0
        )

        # Evaluate the model
        y_pred_scaled_exp = model.predict(X_test_scaled)
        y_pred_exp = scaler_y.inverse_transform(y_pred_scaled_exp)
        y_test_original_exp = scaler_y.inverse_transform(y_test_scaled)

        mae_exp = mean_absolute_error(y_test_original_exp, y_pred_exp)
        mse_exp = mean_squared_error(y_test_original_exp, y_pred_exp)
        rmse_exp = mse_exp ** 0.5
        r2_exp = r2_score(y_test_original_exp, y_pred_exp)
```

```
results.append({
    'optimizer': 'Adam',
    'learning_rate': lr,
    'batch_size': bs,
    'mae': mae_exp,
    'mse': mse_exp,
    'rmse': rmse_exp,
    'r2': r2_exp,
    'history': history.history
})

print(f"  MAE: {mae_exp:.4f}, MSE: {mse_exp:.4f}, R2: {r2_exp:.4f}")

print("\nExperimentation complete. Results stored.")
```

Training with Learning Rate: 0.01, Batch Size: 32

63/63 ━━━━━━ 0s 5ms/step  
MAE: 135.9199, MSE: 57128.4323, R2: 0.6826

Training with Learning Rate: 0.01, Batch Size: 64

63/63 ━━━━ 1s 5ms/step  
MAE: 111.3315, MSE: 51123.6975, R2: 0.7160

Training with Learning Rate: 0.01, Batch Size: 128

63/63 ━━━━ 1s 5ms/step  
MAE: 88.1997, MSE: 42750.5519, R2: 0.7625

Training with Learning Rate: 0.001, Batch Size: 32

63/63 ━━━━ 1s 6ms/step  
MAE: 84.9429, MSE: 36150.7856, R2: 0.7992

Training with Learning Rate: 0.001, Batch Size: 64

63/63 ━━━━ 1s 5ms/step  
MAE: 75.0818, MSE: 33309.9928, R2: 0.8149

Training with Learning Rate: 0.001, Batch Size: 128

63/63 ━━━━ 1s 5ms/step  
MAE: 72.9170, MSE: 32925.7470, R2: 0.8171

Training with Learning Rate: 0.0001, Batch Size: 32

63/63 ━━━━ 1s 5ms/step  
MAE: 85.3650, MSE: 36826.2580, R2: 0.7954

Training with Learning Rate: 0.0001, Batch Size: 64

63/63 ━━━━ 1s 5ms/step  
MAE: 83.0086, MSE: 36822.5719, R2: 0.7954

Training with Learning Rate: 0.0001, Batch Size: 128

63/63 ━━━━ 0s 5ms/step  
MAE: 92.4923, MSE: 39212.0516, R2: 0.7822

Experimentation complete. Results stored.

```
print("\n--- Experiment Results ---")
for res in results:
    print(f"Learning Rate: {res['learning_rate']}, Batch Size: {res['batch_size']}")  
    print(f" MAE: {res['mae']:.4f}")  
    print(f" MSE: {res['mse']:.4f}")  
    print(f" RMSE: {res['rmse']:.4f}")  
    print(f" R2: {res['r2']:.4f}")  
    print("-----")
```

```
--- Experiment Results ---
Learning Rate: 0.01, Batch Size: 32
MAE: 102.6197
MSE: 48909.6911
RMSE: 221.1554
R2: 0.7283
```

```
-----  
Learning Rate: 0.01, Batch Size: 64
MAE: 89.3748
MSE: 39402.5446
RMSE: 198.5007
R2: 0.7811
```

```
-----  
Learning Rate: 0.01, Batch Size: 128
MAE: 100.0023
MSE: 44490.8832
RMSE: 210.9286
R2: 0.7528
```

```
-----  
Learning Rate: 0.001, Batch Size: 32
MAE: 76.9474
MSE: 34480.3880
RMSE: 185.6890
R2: 0.8084
```

```
-----  
Learning Rate: 0.001, Batch Size: 64
MAE: 79.7325
MSE: 38531.9798
RMSE: 196.2956
R2: 0.7859
```

```
-----  
Learning Rate: 0.001, Batch Size: 128
MAE: 75.6608
MSE: 32865.7719
RMSE: 181.2892
R2: 0.8174
```

```
-----  
Learning Rate: 0.0001, Batch Size: 32
MAE: 80.1344
MSE: 35087.2408
RMSE: 187.3159
R2: 0.8051
```